

APPLICATION CONFIGURATION DATA SERVICE

Author: Rijn Buve – Reviewers: Charles Davies, Kees van Boxel, Kees Schuerman, Andreas Wuest, Sven Baselau

Version: 2016-09-26

Introduction	1
Functional Requirements	1
Non-Functional Requirements	2
Service Description	2
API Description	3
Matching the Requirements	9

Introduction

This document describes the requirements for and API of a generic service to read key-value pairs from a server, given a number of criteria. The key-value are referred to as “parameters” and generally represent application configuration data for the client.¹

Typical use cases for this service are:

- Application or system configuration of clients
- Service discovery for clients

Typical operational contexts for this services are:

- Cases where many clients shared the same configuration but you wish to have the ability to specify special cases (e.g. for testing)
- Clients may wish to apply controls to reduce data usage (such as for over-the-air data plans) by using configurable service parameters (such as request frequency or scope)
- You wish to have controls to minimize the amount of duplication between mostly shared but slightly differing configurations
- High demand, high availability environments
- The ability to apply robustness controls, such as safe fallback configurations

The service is essentially defined as a read-only key-value store. It returns string values for keys by default, but it is perfectly fine for a client to provide its configuration as a binary “blob” (represented in the value of 1 key-value pair).

The service is generic and can potentially be used in many environments, such as environments for automotive clients (for CS), mobile apps, PNDs or even back-end web services.

The service is not specifically designed to be multi-tenanted (although it can be used as such). As a micro-service, it makes sense to deploy an instance of the service per context. It has been designed to be easily and highly scalable and to be able to run at low resource cost.

Functional Requirements

We define a “configuration” as a set of key/value-pairs which are context specific (i.e. the client knows how to interpret the key/values, not the service).

We recognize the following functional requirements for the parameter service:

¹ This service may be considered a “micro-service”: a narrowly focused, independently deployable service.

1. The service offers a REST API to fetch a configuration based on client-provided hierarchical criteria.
2. The service provides a way to reconfigure the system.
3. The service offers the ability to provide string values for key-value pairs. The strings may represent a string, integer, or any other (possibly encoded) type (such as base64 encoded binary blobs).
4. The service offers controls to minimize the number of configurations that need to be stored for clients.
5. The service offers controls to minimize data usage by clients.
6. The service supports JSON or XML (or both) for the specification and retrieval of configurations.

Non-Functional Requirements

We recognize the following non-functional, or quality requirements for the parameter service:

1. The service needs to support linear horizontal scalability, both in performance and availability.
2. The service needs to support a sufficient workload capacity; P98² >= 50 concurrent requests.
3. The service needs to support a sufficient throughput; P98 <= 50 msec.
4. The service needs to offer high-availability; 99.9% per instance.
5. The service needs to “run cheaply”; must run on an AWS “t2.medium” instance or less (given a reasonable set of configurations).
6. The service is generic in a way that there no references to device-related, automotive-related or context-specific concepts.
7. The service (as a whole) does not require downtime when reconfiguring. (Individual nodes may go down, if the service is still able to operate well.)
8. The service requires no authentication/authorization (to be dealt with outside this service).

Service Description

The service can be visualized as operating on a search tree, which consists of nodes and leafs. The nodes themselves represent hierarchical search terms; the leafs are the search results, see Figure 1.

Given a set of hierarchical search terms, the service will try to match the search terms in the search tree and returns the leaf of the deepest node still matching the terms.

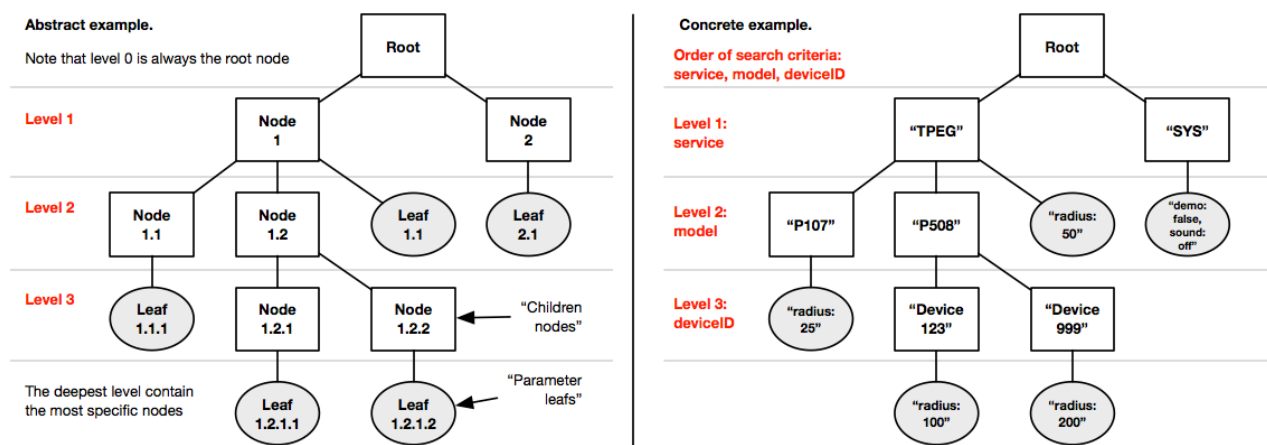


Figure 1. Search tree with nodes (search terms) and leafs (search results).

Example. Consider the example tree on the right in Figure 1.

This example represents a service configuration in which devices wish to be able to get their configuration parameters separately for TPEG services and system services. Every device has a “model type”, which in this case is “P107” or “P508” and a unique device ID.

² P98 stands for the “98th percentile” of a population (only 2% of the population does not fall in this category).

The square nodes in the tree represent the search terms that need to be matched and the round leaves contain the search results.³

The service should produce the following search results, if a client provides its search criteria:

- service=**TPEG**, model=**P508**, deviceId=**Device123**, returns **radius:100**, returns leaf of the **Device123** node.
- service=**TPEG**, model=**P508**, deviceId=**Device456**, returns **radius:50**, falls back to the leaf of the **TPEG** node.
- service=**TPEG**, model=**P107**, deviceId=**Device789**, returns **radius:25**, falls back to the leaf of the **P107** node.
- service=**SYS**, model=**P508**, deviceId=**Device123**, returns **demo:false**, falls back to leaf of **SYS** node.

If no matching nodes are found, or no leaf is attached to highest matching node, a “no match found” code is returned by the service (which in HTTP would be represented by a 404 status code).

The service allows the integrator to specify an initial search tree, which is read from a configuration file as well as an integration REST API to manage the search tree.

API Description

JSON or XML Configuration Files and Responses

The service accepts both JSON and XML as the input for configuration files and the output for responses. The default response format for the REST API is JSON (if the HTTP “**Accept:**” header is omitted or is “**Accept:application/json**”). To retrieve XML response, simply specify “**Accept:application/xml**”.

Note that this document uses JSON for most of its examples, but every configuration files or response is allowed to be both XML and JSON.

Initializing the Search Tree

The service provides the ability to specify a URI property for the service, from which the initial configuration of the search tree is retrieved. We call this URI the **startupConfigurationURI** and it is specified in a properties file (or by other means of injection). The property must be available during start-up or the service will not boot (and produce a human understandable error entry in the log as to why it will not boot).

During start-up, the **startupConfigurationURI** is read, which may be either a URL (prefixed “http[s]:”), a file from a shared drive (prefixed “file:”), or a file on the classpath (prefixed “classpath:”). The format of the configuration file can be either JSON or XML. The service will figure out which one it is.

The contents are interpreted as the contents of the root node of the search tree. If the URI was not specified, it could not be found, or its format was incorrect, an error entry will be put in the log and any subsequent call to retrieve information about the tree will return a HTTP FORBIDDEN (403) status code.

Monitoring hooks should represent this error state as well. (For example, if the service is to be deployed on Amazon Elastic Beanstalk with an HTTP “/status” monitoring call and initialization fails, the status call returns HTTP FORBIDDEN (403)).

The service is expected to be run as a bunch of parallel nodes with the same configuration, connected to a load balancer. When a change to the search tree is required, you should edit the configuration file, or redirect the properties file to another configuration and restart all instances/nodes of the service, one-by-one (see also next paragraph).

Adding additional server nodes is as easy as attaching more nodes to the load balancer and booting them.

If you wish to retrieve the configuration of a node, simply execute “GET /tree”. The response matches the original configuration tree.

³ Note that although in this example only a single parameter is returned (radius=xyz), you would normally specify any number of parameters as a search result.

Providing a New Configuration to Immutable Service Nodes

After a service node has read its configuration, it becomes immutable and cannot be altered in any way. The service provides no API to modify the configuration of nodes during runtime other than stopping nodes and restarting them with a different configuration which will be read during start-up.

This approach may seem inflexible, but is actually extremely simple, reliable, robust and removes the need for intra-node communication, shared databases and such. All you need to do to reconfigure a running system, is provide a new configuration file and restart each node, one by one, which will cause zero downtime for the end-user. (And a service like Amazon Elastic Beanstalk will do this automatically for you.)

Querying Specific Nodes from the Search Tree

The service exposes REST methods to query the configuration of the search tree. (Don't confuse the querying the configuration of the search tree with searching the search tree for results.)

```
GET /tree[/level1[/level2[...]] -- get specific tree nodes
```

This produces a response like:

```
{
  "nodes": ["{node1}", "{node2}", ...],
  "parameters": [{"key": "{key1}"}, {"value": "{value1}"}, {...}, ...]
}
```

The "nodes" array is optional and lists the children nodes with search terms one level below the specified node. The "parameters" value is the optional leaf node of this node and lists the search result (an array of key-value pairs).

Important: Normally, you would not use this method to get a configuration for a device, as it requires you to *exactly* specify a node in a hierarchy and does not evaluate any regular expression. (You would typically use this only in a program that wants to read the actual nodes, for example, for display purposes).

Normally, you would use the "**search**" method below, because it provides fallback configurations if certain elements of the search path cannot be matched. The method to fetch individual nodes is usually only useful for the root node, to fetch the entire configuration from a node.

Search for the Best Matching Parameters

This is the method clients would normally use to fetch a configuration. It is used to traverse the search tree for the best matching search result. The returned result may either be an exact match for the input, or be the best "fallback" configuration. Its format is:

```
GET /tree?order={level1}/{level2}/...&search={value1}/{value2}/...
```

The 'order' element specifies the order in which the node level names will be specified in the search. The 'search' element specifies the values of the elements (in the order just specified). String matches are always case-sensitive.

For example, if the configuration tree has two levels, where the first level nodes represent the "serviceID" and the second level nodes represent the "deviceID", then the client may specify a search either as:

```
GET levels=serviceID/clientID&search=someService/someDevice
```

or as

```
GET levels=clientID/serviceID&search=someDevice/someService
```

Important: Note how the order of the levels is redefined in the search query. This means clients do not need to be changed whenever you wish to reshuffle the node levels of the configuration tree at the server, to reduce its complexity. (To know which levels are available, simply retrieve the top-level tree, using "GET /tree", which contains the element "levels" with the order of level names in the tree.)

The returned result contains the parameters of the leaf of the deepest node matching the search criteria as well as the path of the node to which the leaf is attached:

```
{ "parameters": [{"key": "{key1}", "value": "{value1}"}, {...}, ...],
  "matched": "..." }
```

If no search result was found, HTTP status code 404 (NOT FOUND) is returned instead. The **"matched"** attribute is particularly useful to understand whether fallback parameters were used and if so, from which node. For example, if a query was executed for **"x=1,y=2,z=3"**, but the configuration did not specify a value for a node with z=3, then the returned path could indicate the actual node it returned would be **"x=1,y=2"**, rather than **"x=1,y=2,z=3"**.

Note that the fallback mechanism is entirely unambiguous and deterministic due to the hierarchical nature of the search tree configuration. This means fallbacks should never need to cause unexpected surprises.

Example. Consider the example tree on the right in Figure 1 again.

Using the search method, you would get these results:

```
GET /tree?levels=service/model/deviceID&search=SYS
=> demo=false, sound=off (exact match)
GET /tree?levels=service/model/deviceID&search=SYS, P508, Device123
=> demo=false, sound=off (fallback to SYS)
GET /tree?levels=service/model/deviceID&search=TPEG
=> radius=50 (exact match)
GET /tree?levels=service/model/deviceID&search=TPEG/P508/Device123
=> radius=100 (exact match)
GET /tree?levels=service/model/deviceID&search=TPEG/P508/Device456
=> radius=50 (fallback to P508)
GET /tree?levels=service/model/deviceID&search=TPEG/P107/Device789
=> radius=25 (fallback to P107)
GET /tree?levels=service/model/deviceID&search=OTHER
=> NOT FOUND (no fallback found)
```

In contrast, using the method to query the actual tree nodes (which you normally would not use), you would get these results:

```
GET / => produces help text for service
GET /tree => levels=service,model,deviceID, nodes={...}
GET /tree/SYS => parameters: demo=false, sound=off
GET /tree/SYS/P508/Device123 => NOT FOUND
GET /tree/TPEG => nodes: P107, P508; parameters: radius=50
GET /tree/TPEG/P508/Device123 => parameters: radius=100
GET /tree/TPEG/P508/Device456 => NOT FOUND
GET /tree/TPEG/P107/Device789 => NOT FOUND
```

HTTP Status Codes

The HTTP status codes for the service are defined as:

- Successful GET: 200 – OK
- Node not found for GET: 404 – NOT FOUND
- No search result found for GET: 404 – NOT FOUND

Specifying the Configuration File

The configuration file is read from the `"ApplicationConfigurationData.startupConfigurationURI"` property in the properties file of the service. Below are some examples for specifying configuration files for nodes

to read during start-up. We also provide 2 examples of the configuration file to produce the configuration tree of Figure 1 (one as JSON and one as XML).

Examples:

Specify a URL location for an initial configuration (assuming a properties file):

```
ApplicationConfigurationData.startupConfigurationURI = http://datacenter/myconf.json
```

Specify a configuration file on a shared drive:

```
ApplicationConfigurationData.startupConfigurationURI = file://share-1/myconf.xml
```

Specify a configuration file on the classpath:

```
ApplicationConfigurationData.startupConfigurationURI = classpath://myconf.xml
```

Configuration file for the search tree in Figure 1 in JSON format.

```
{
  "levels": ["service", "model", "deviceID"],
  "nodes": [
    {
      "name": "TPEG",
      "nodes": [
        {
          "name": "P107",
          "parameters": [{"key": "radius", "value": "25"}]
        }, {
          "name": "P508",
          "nodes": [
            {
              "name": "Device123",
              "parameters": [{"key": "radius", "value": "100"}]
            }, {
              "name": "Device999",
              "parameters": [{"key": "radius", "value": "200"}]
            }
          ]
        }
      ]
    }, {
      "name": "SYS",
      "parameters": [{"key": "demo", "value": "false"}, {"key": "sound", "value": "off"}]
    }
  ],
  "modified": "2016-04-05T17:28:16Z"
}
```

Configuration file for the search tree in Figure 1 in XML format.

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<tree>
  <levels>
    <level>service</level>
    <level>model</level>
    <level>deviceID</level>
  </levels>
  <nodes>
    <node>
      <name>TPEG</name>
      <nodes>
        <node>
          <name>P107</name>
          <parameters>
            <parameter>
              <key>radius</key>
              <value>25</value>
            </parameter>
          </parameters>
        </node>
        <node>
          <name>P508</name>
          <nodes>
            <node>
```

```

        <name>Device123</name>
        <parameters>
            <parameter>
                <key>radius</key>
                <value>100</value>
            </parameter>
        </parameters>
    </node>
    <node>
        <name>Device999</name>
        <parameters>
            <parameter>
                <key>radius</key>
                <value>200</value>
            </parameter>
        </parameters>
    </node>
</nodes>
</node>
</nodes>
<parameters>
    <parameter>
        <key>radius</key>
        <value>50</value>
    </parameter>
</parameters>
</node>
<node>
    <name>SYS</name>
    <parameters>
        <parameter>
            <key>demo</key>
            <value>false</value>
        </parameter>
        <parameter>
            <key>sound</key>
            <value>off</value>
        </parameter>
    </parameters>
</node>
</nodes>
<modified>2016-04-05T17:28:16Z</modified>
</tree>

```

Using Regex's to Match Ranges of Criteria

The service allows you to specify regex's in the configuration "name" values. In that case, specified names in a search query are matched against these regex's. This allows you to specify configuration for multiple search paths (read: devices) in a simple way.

For example, if the configuration contains a node:

```
"name": "Device100[0-9]+", "nodes": [ ... ], ...
```

Then, all of the queries below will be served the same parameters:

```

GET /tree?levels=device&search=Device1002
GET /tree?levels=device&search=Device10072
GET /tree?levels=device&search=Device100846

```

This makes it much easier to specify exceptional configurations, for example, for test devices.

Note that "exact" string matches always prevail over regex matches. So, if a search string "Joe" is provided, a match which node name "Joe" always prevails over node name "J.*".

Reducing Data Usage Using HTTP Header If-Modified-Since

To minimize data usage, specifically over-the-air, the service allows the use of the standard HTTP header "If-Modified-Since" or "If-None-Match" (next paragraph). The supplied value is matched to the "modified" date of the root node (see the example above) and the service will return HTTP status code 304 (NOT MODIFIED) if the root node is not newer than the supplied date/time.

This means the “modified” attribute can only be supplied for the root node. The modified date/time is valid for the entire configuration file. Its format is:

```
"modified": "1994-11-06T07:49:37Z"
```

Note that the format of the “If-Modified-Since” HTTP header is defined by W3C, RFC 1123, as:

```
If-Modified-Since: Sun, 06 Nov 1994 08:49:37 GMT
```

Failure to provide the correct format for the HTTP header (or accidentally use the ISO format) results in ignoring the header, which means a full result is always returned.

Reducing Data Usage Using HTTP Header If-None-Match

You can also use “ETags” to reduce data consumption. ETags may be thought of as the hash of the response: if two response are the same, they have the same ETag and if the ETags differs, the responses differ.

The service returns an ETag value with each request and the caller may supply that ETag in the HTTP header “If-None-Match” header at a next request. If the response since last time is the same, the service will return HTTP status code 304 (NOT MODIFIED) rather than providing the full response again.

The ETag is provided by the service in the HTTP as:

```
ETag: "686897696a7c876b7e"
```

And the caller may provide this ETag in a subsequent request as:

```
If-None-Match: "686897696a7c876b7e"
```

Failure to provide the correct format for the HTTP header (or accidentally use the ISO format) results in ignoring the header, which means a full result is always returned.

Important: Which method you should choose, If-Modified-Since or If-None-Match, is up to you. The If-None-Match may be a better choice if you expect regular changes to the configuration tree, which affect only a subset of the callers (such as adding test device IDs).

Note also that if the supplied ETag does not match, a full response is always returned, regardless of the modified date (because apparently you got a different response earlier).

Sharing Configuration Subtrees with Includes

It is possible to share subtrees of the configuration tree for different configurations. For example, if a specific subtree resembles the settings for a service and these settings are shared for many device configuration, you may wish to specify these settings once in a separate configuration URI and simply include that URI anywhere in the configuration tree.

This looks like this:

```
{
  "nodes": [
    { "include" : "file://shared/sys_config.json" },
    { "include" : "file://shared/service1_config.json" },
    { "include" : "file://shared/service2_config.json" }
  ],
  "modified" : "2016-04-05T17:28:16Z"
}
```

The include file "sys_config.json" would live on the class path in this case and might look like this:

```
{
  "name": "SYS",
  "parameters": [{"key": "demo", "value": "false"}, {"key": "sound", "value": "off"}]
}
```

Include files can live on the classpath (classpath:), on a shared drive (file:) or web-service (http:/https:). Nodes names from peer nodes must be unique, also when include files are used, or the service won't start.

Include files can be nested to any level, but cannot be included recursively (as the recursion would not end). This allows you, for example, to create a:

- config file per high-level service, which includes
 - config files per device configuration, which include
 - config files for individual devices

to create a hierarchy of config files, which may be easier to maintain than 1 enormous config file.

Note that only the “modified” time stamp of the root node is considered for the HTTP If-Modified-Since optimization discussed earlier.

Getting Multiple Configurations at Once

It is possible to query the search tree for multiple results at once, reducing the number of calls a client needs to make to the service. Separate searches are separated by a single ‘;’ in the search parameters (and individual search terms are separated by ‘/’ as usual):

This looks like:

```
GET /tree?levels=name1/name2/...&search=value1/value2,...;valueX/valueY,...
```

Note the individual searches are separated by a semi-colon. The returned response is in that case a JSON array of results, rather than a single result. The elements of the array are in the same order as the search query elements.

Note that if any of the sub-queries are not found, the entire query responds with 404 (NOT FOUND), (rather than supplying an empty response item, which could be misleading).

Matching the Requirements

The service as defined in this document has been implemented at

<https://bitbucket.tomtomgroup.com/projects/ST/repos/application-configuration-data-service/browse>

It matches the requirements listed as follows:

- Reconfiguring the system is done by providing a new configuration setup and rebooting the nodes one-by-one. This will cause zero downtime of the system as whole and is extremely reliable as a mechanism.
- The number of configurations can be kept low by using the “include” mechanism wisely and providing smart fallback configurations, or using regular expressions in the configuration tree to have one configuration apply to many clients.
- The service is stateless, which makes it trivial to up- or downscale the nodes horizontally.
- The service does not require a database or other middleware, which makes it extremely simple to deploy, monitor and operate. This also reduces operational costs to a minimum (requiring the smallest of virtual machines).
- Everything is served from memory, which produces extremely fast responses.

Change Log

2015-08-05 Initial version, extracted from document “Perseus Integration for CS”. Corrected key-value type to array type.

2015-08-07 Removed reference to JSON, improved HTTP status code results.

2015-08-10 Added the ability to read an initial configuration setup from a (properties specified) URI. Also added management interface for changing parameters of a node only (POST).

2015-08-12 Checked final comments.

2015-09-17 Removed remaining comments.

2016-04-05 Removed all write/update methods. The service is now defined to be read-only. Added If-Modified-Since.

2016-04-11 Added include functionality.

2016-07-25 Updated document to include full XML support and corrected examples.

2016-09-22 Major change to API: GET no longer hierarchically structured, only the configuration tree is.

2016-09-26 Renamed parameter 'order' to 'levels' for consistency.